

Project Proposal: skyjam, Forecasting GNSS Interference over European Airspace

Zayden (github.com/Zayden16) · MLOps HS26 · MS1 Proposal
Repository (public): <https://github.com/Zayden16/hslu-mlops-hs-26>

1 Problem statement

What and for whom. GNSS jamming and spoofing are now routine over the Baltic, the Black Sea and the Eastern Mediterranean: aircraft lose the integrity of their satellite position fix mid-flight. Public tools such as gpsjam.org show *yesterday's* map. A dispatcher planning a route, a drone operator or a flight-ops centre needs the opposite: a statement about the hours ahead.

Prediction. For every H3 resolution-2 airspace cell (~180 km edge), predict the probability that the cell is *interference-affected exactly 6 hours ahead*. This is binary classification per cell-hour, re-issued every hour. **Scope:** European and adjacent airspace, covered by 18 sampling discs of 250 NM (11 over regions with reported interference, 7 over Western-European control regions), about 80 populated cells per sweep.

Label. A cell-hour is affected when at least **30 %** of the cruise-altitude aircraft in it broadcast a degraded navigation integrity category (NIC < 7, containment radius worse than 0.2 NM). NIC is each aircraft's own assessment of its position quality, sent over ADS-B. The label is a *fleet-level consensus about a physical event* at a future timestamp, aggregated over independent airframes, not a rule applied to the model's inputs.

Success criterion. On a strictly out-of-time test split the model must beat two baselines:

Baseline	Why it is the bar
Persistence: state at t repeated at $t + 6$ h	Interference is strongly autocorrelated; not beating this means nothing was learned
Per-cell hourly climatology	Encodes “the Gulf of Finland is usually bad at 14:00” with no live signal

Target: PR-AUC at least 0.10 above persistence and a lower Brier score, with the gain holding *separately* in interference and control regions. **Rare positive class** (4.5 % of cells in the first validation snapshot): accuracy is meaningless, so PR-AUC is the headline metric. It is handled with class weighting and a decision threshold tuned on the validation split, and precision/recall are reported per region, because a model that permanently flags the whole Baltic would look good globally and be operationally useless.

2 Originality & motivation

Checked against the HSLU archive (mlops-lab.ch) and the KTH ID2223 project lists: both are dominated by crypto prices, weather, air quality and transport availability. None predicts the *quality* of a navigation signal, and no ADS-B project appears. **In one sentence: the dataset does not exist until my pipeline creates it, and the label is inferred from fleet telemetry rather than read from a column.**

- No historical download exists** for per-cell navigation integrity. The feature pipeline is the sole source of history, so the scheduled job is load-bearing, not decorative.
- The label is constructed** from a fleet-level statistic over self-reported avionics quality, with a defensible threshold and an explicit traffic floor.
- Drift is the subject, not an afterthought.** A new emitter switching on is genuine concept drift in its region, so MS4 monitoring measures the very phenomenon being modelled.

Motivation: GNSS interference is an active civil-aviation safety concern in Europe, the raw telemetry is fully public, and I want a project whose failure modes are physical rather than financial.

3 Data source & features

Source. `api.adsb.lol` v2, a public community ADS-B aggregator, endpoint `/v2/point/{lat}/{lon}/{radius}`. A documented JSON API, **not scraping**, free and key-less for non-commercial use. **Fallback:** `api.adsb.fi` serves the same schema and is a one-variable switch (`SKYJAM_ADSB_BASE_URL`).

Update frequency and volume. The feed is real time. The pipeline sweeps all 18 discs every **30 minutes**, pacing requests 12 s apart (the measured rate limit) with exponential backoff. History **starts when my poller starts**; there is no back-download. As of 26.09.2026 the store holds **1.09 M observations over 189 hourly slots** and grows by roughly 140 k observations, or about 1 900 labelled cell-hours, per day, so MS3 will train on the order of 100 000 cell-hours.

Validated signal (live snapshot, aircraft at or above FL200, degraded = NIC < 7):

Interference region	Deg.	Interference region	Deg.	Control region	Deg.
Gulf of Finland	52.0 %	Baltic south	13.2 %	Germany	2.4 %
Poland east	34.0 %	Romania/Moldova	6.2 %	Switzerland	1.4 %
Kaliningrad	25.9 %	Black Sea west	5.1 %	UK / France	1.0 / 0.0 %

A roughly 40-fold separation between worst region and controls: the target is real, not assumed. **The altitude filter matters:** unfiltered, the Swiss disc shows 8.9 % degraded aircraft against 2.3 % at FL200+, because cheap general-aviation avionics report poor integrity everywhere; the same filter leaves the Baltic figures essentially unchanged.

Named features per cell and hourly slot: `aircraft_count`, `degraded_share`, `median_nic`, `p25_nic`, `mean_nac_p`; lags of `degraded_share` and of the binary state at 1, 2, 3, 6, 12, 24 and 168 h; backward-looking rolling means of `degraded_share` over 6 h and 24 h; `aircraft_count_mean_24h`; calendar features `hour_utc`, `day_of_week`.

Leakage. (i) *Split*: strictly chronological train / validation / test with a 6-hour gap between segments so no example straddles a boundary; no shuffling anywhere. (ii) *Temporal, not positional lags*: each cell is reindexed onto a gap-free hourly grid first, because a positional `shift(1)` would present a three-hour-old value as one hour old whenever a sweep was missed; a unit test asserts the lag across a removed hour is NaN. (iii) *Rolling windows exclude the present* (shifted one slot, tested). (iv) *Unobserved is not quiet*: cells under five aircraft are dropped, not imputed as clean, and rows whose future state was never observed carry no label.

4 System design




Figure 1: FTI architecture. Blue: the three pipelines; orange: stores; solid arrows: data flow; dashed: retrain trigger. Deployed today: the live source, feature pipeline and raw store.

Pipelines and triggers. (1) *Feature*: a Go service on Railway sweeps the 18 discs every 30 minutes and writes **unfiltered** observations to PostgreSQL: raw coordinates, nullable altitude, no H3 cell and no label. Thresholds are applied at *read* time, so the altitude floor or H3 resolution can change and the whole history be rebuilt. (2) *Training*: weekly and on drift; reads the store, splits chronologically, trains LightGBM against both baselines, logs to MLflow, registers the best run. (3) *Inference*: hourly and on UI request; loads the Production model and writes a 6-hour forecast per cell.

Tech stack, one line of rationale each:

- **Cloud and feature store**: Railway (container hosting) + PostgreSQL. Durable, unlike the 90-day CI artifacts it replaced, and SQL applies thresholds at read time.
- **Orchestration**: a long-lived Go service with a wall-clock ticker. Four staggered GitHub Actions crons captured only 53 % of hours over 74 h (GitHub drops scheduled triggers and never retries), and here a missed hour is lost forever.
- **Tracking and registry**: MLflow, for run comparison and a Production alias that inference loads, so the deployed version is always identifiable.
- **Serving**: Streamlit map on Railway, next to the data; a map is the natural view of a spatial forecast.
- **Packaging and CI**: Docker, `uv.lock` for Python, Go modules; GitHub Actions runs lint and 50 tests (34 Python, 16 Go) against a real PostgreSQL.

Training-serving skew is prevented structurally: thresholds, altitude floor, H3 resolution and lag set live once in `src/skyjam/common/schema.py` and are applied on read. The Go capture service computes **no** label, and the sampling grid it shares with Python has a parity test that fails the build if the two drift.

Optional / stretch (not core): Terraform IaC, Airflow, a managed feature store, alerting on the drift metric. Core FTI comes first.

Current state and risks. The capture service is deployed and has completed 18/18 discs on every recent sweep. Risks: history accrues only from now on (mitigated by capturing since proposal time and storing raw, so a definition change does not invalidate it); the API could throttle or disappear (measured pacing, backoff, `adsb.fi` fallback); sparse night traffic (per-cell traffic floor, unobserved cells excluded). **The repository is public**: <https://github.com/Zayden16/hslu-mlops-hs-26>.